



# Intro Into Index Analysis

Dmitri Korotkevitch  
<https://aboutsqlserver.com>

# Hello!

20+ years of experience in IT

15+ years of experience working with SQL Server

Microsoft Data Platform MVP

Microsoft Certified Master

Author




- Advanced SQL Server Troubleshooting and Performance Tuning
- Pro SQL Server Internals (v1-2)
- Expert SQL Server In-Memory OLTP (v1-2)
- Expert SQL Server Transactions and Locking

Blog: <https://aboutsqlserver.com>

Email: [dk@aboutsqlserver.com](mailto:dk@aboutsqlserver.com)

Slides / Code: <https://github.com/aboutsqlserver/code>

YouTube: [AboutSQLServer](#)



# Agenda

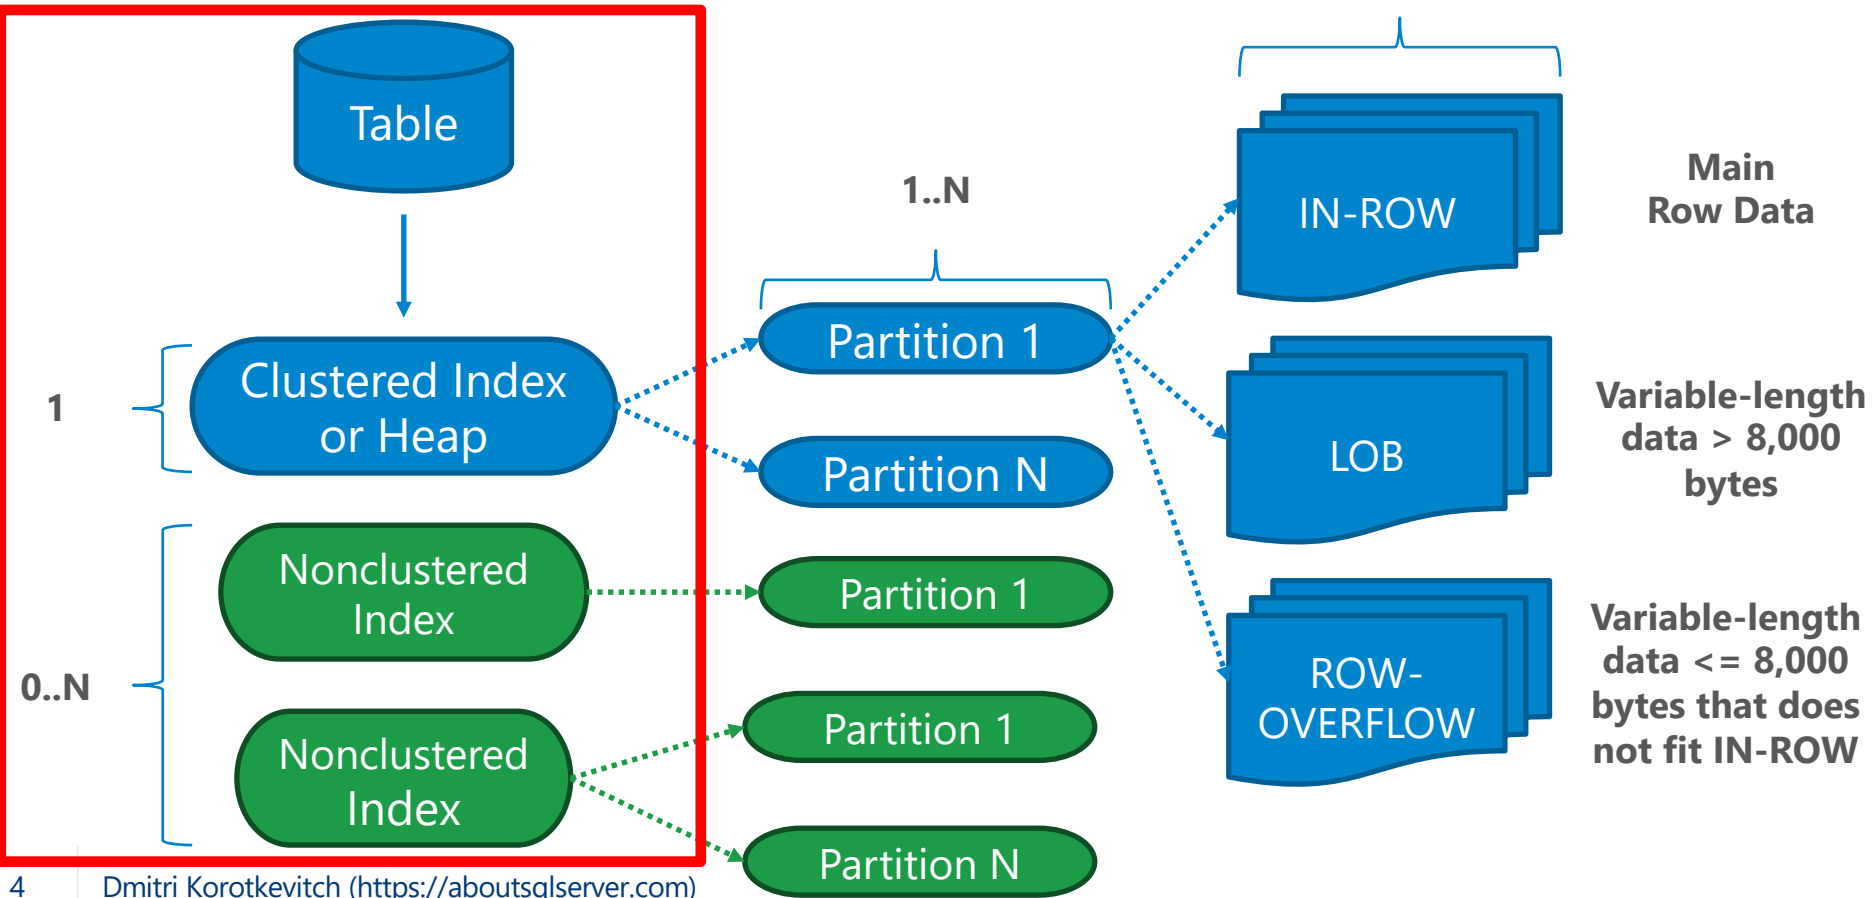
## Data Storage 101

- B-Tree Format
- Access Patterns


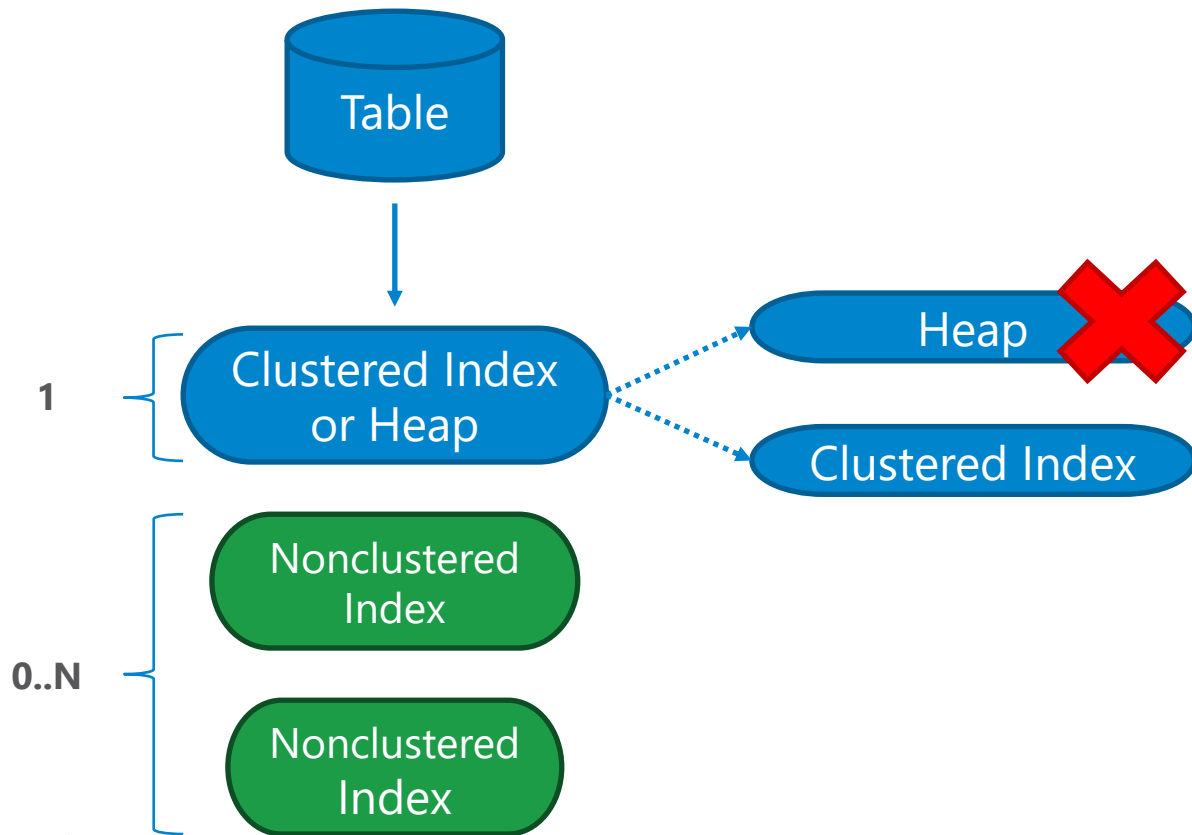
## Index Analysis Techniques

- Catalog views
- Usage and Operational statistics
- Physical statistics

# Data Storage 101



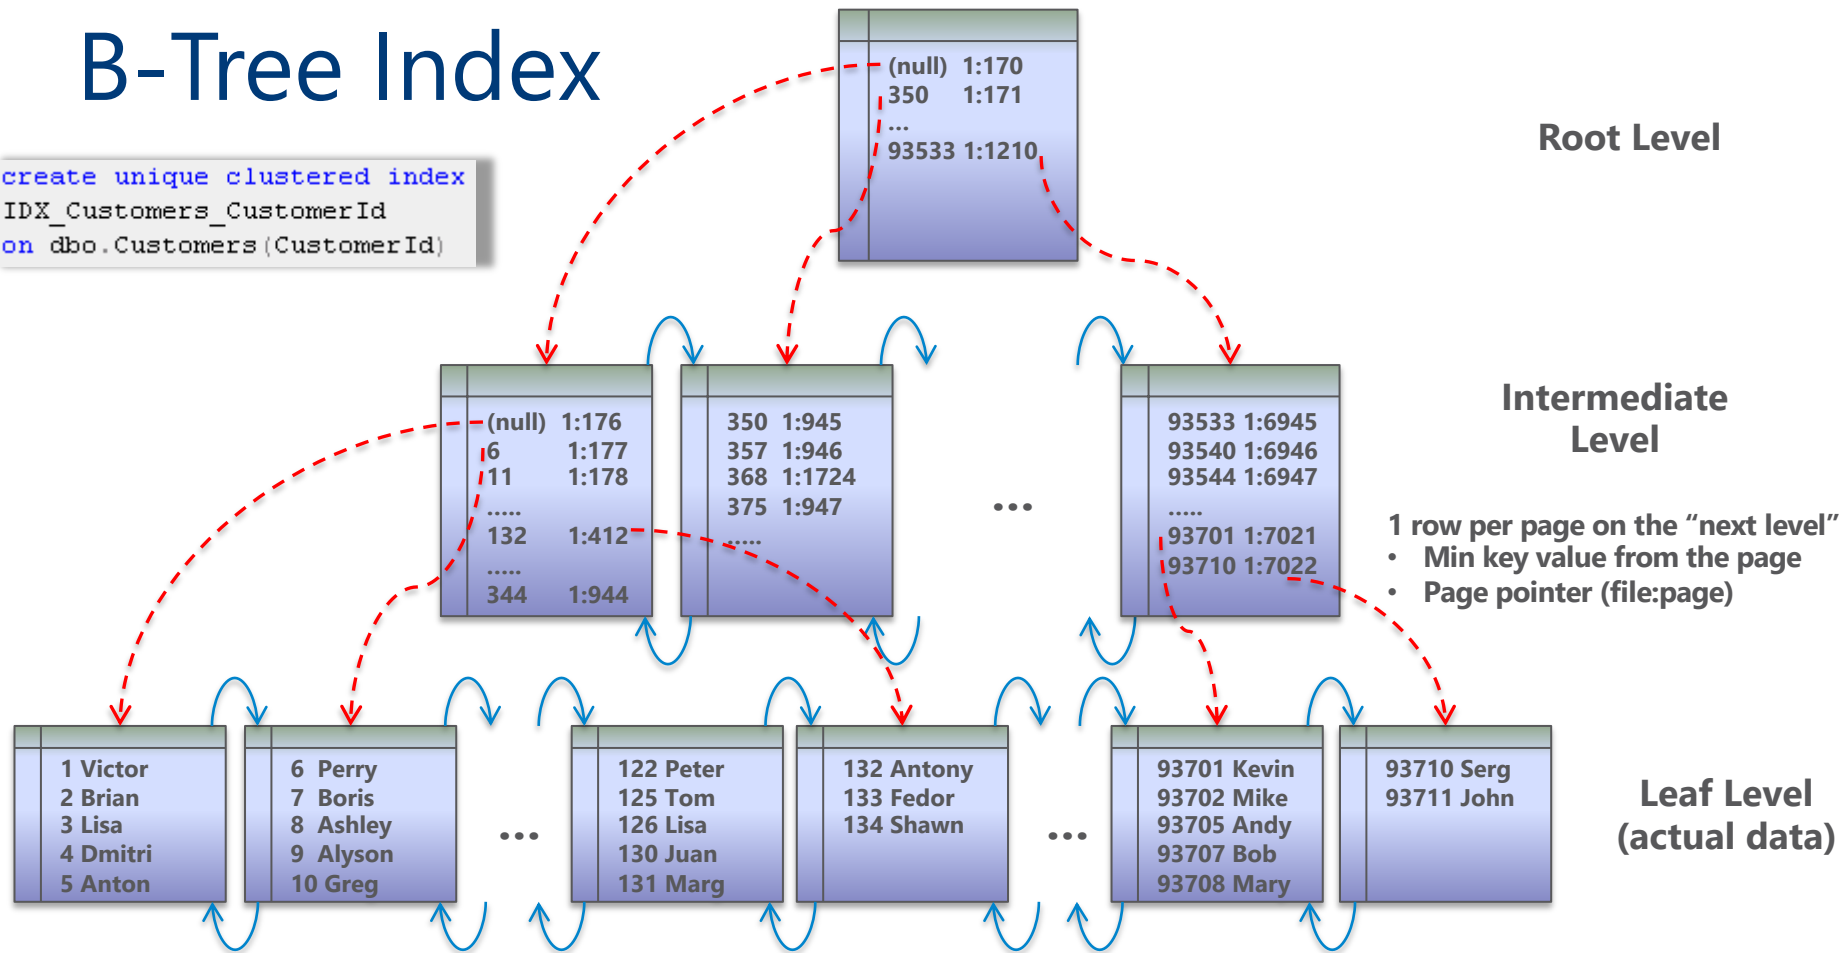
# Data Storage 101



<https://www.flickr.com/photos/yorokobi/473323374>

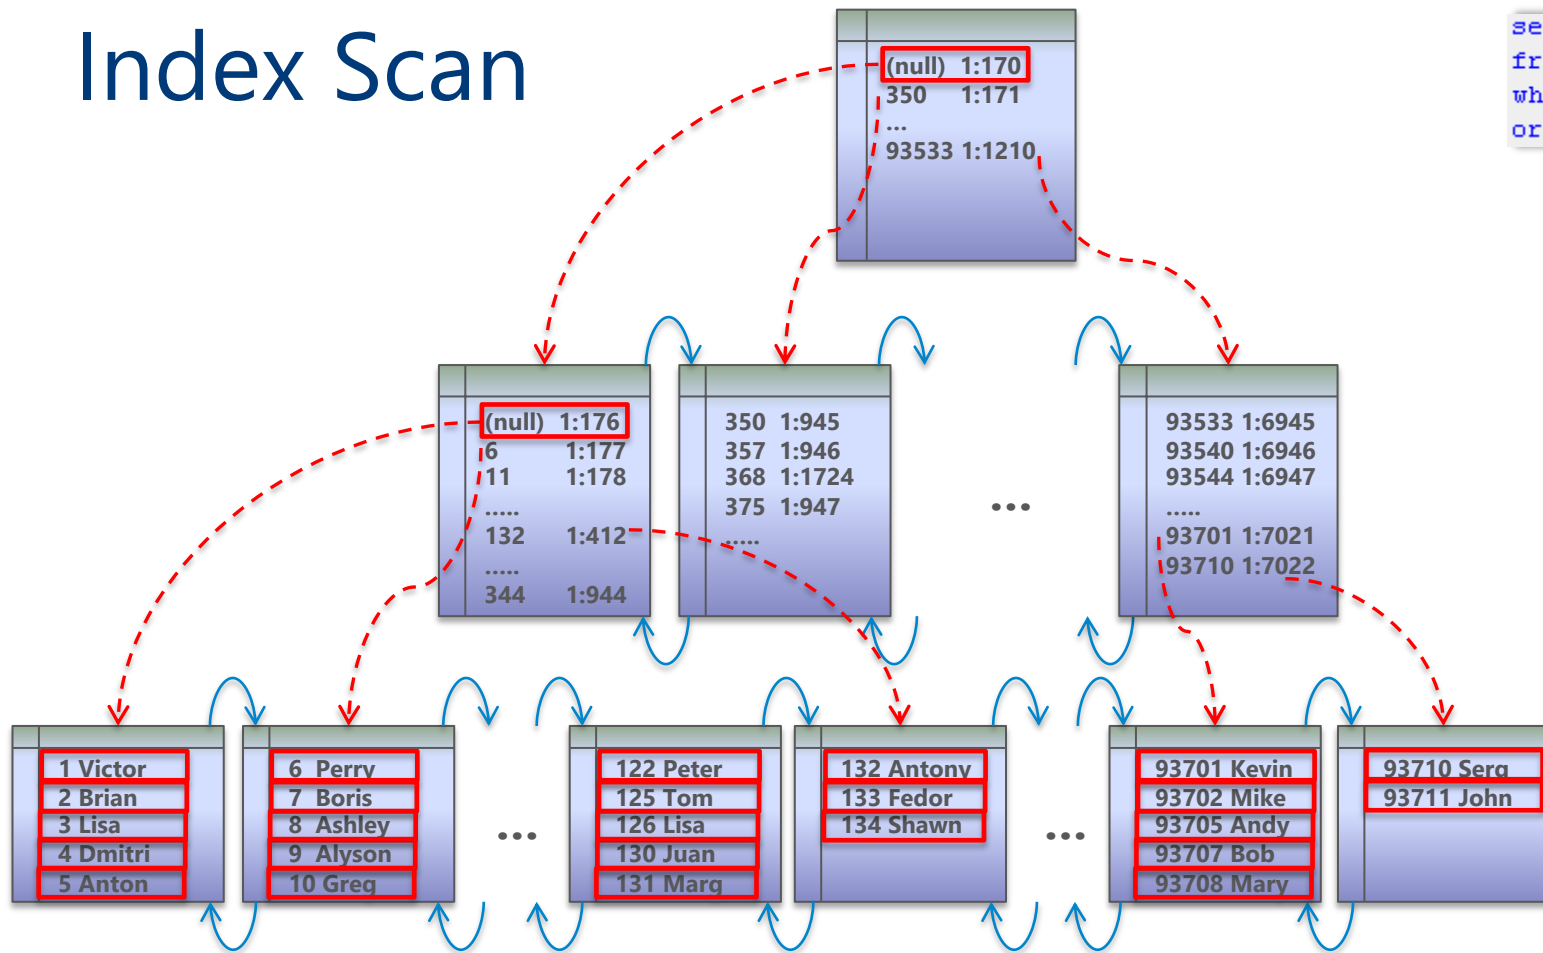
# B-Tree Index

```
create unique clustered index  
IDX_Customers_CustomerId  
on dbo.Customers (CustomerId)
```



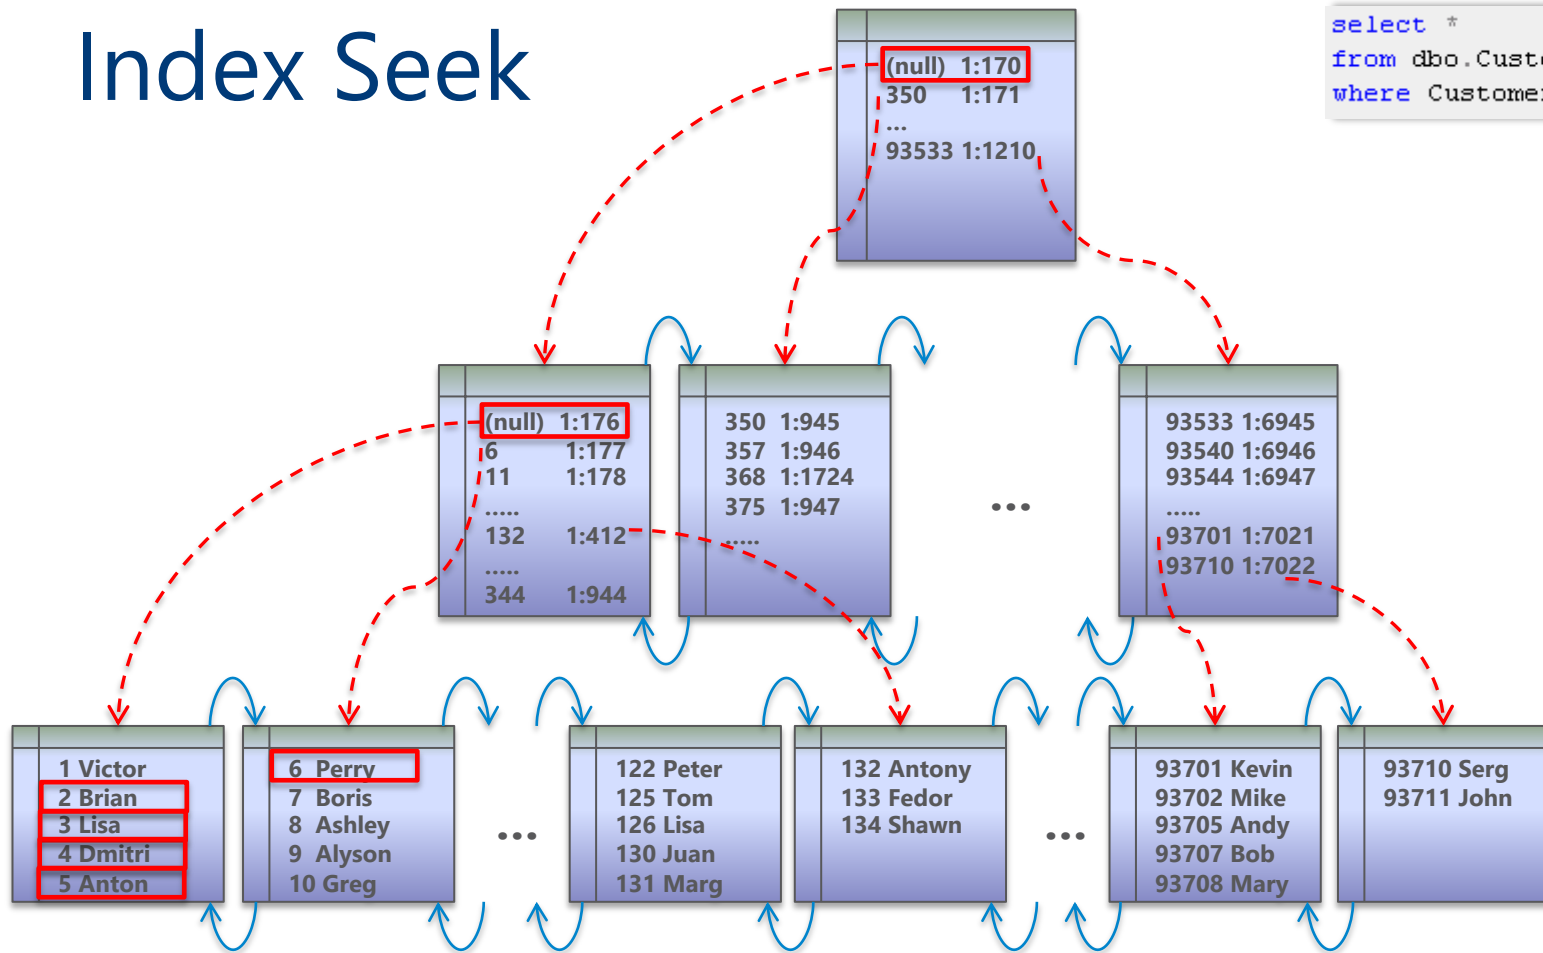
# Index Scan

```
select *  
from dbo.Customers  
where Name = N'Boris'  
order by CustomerId
```



# Index Seek

```
select *  
from dbo.Customers  
where CustomerId between 2 and 6
```







# Seeks and Scans

## Demo

# Patterns to Avoid

## Functions

### Substring search

- Prefix search is OK

### Data type conversion

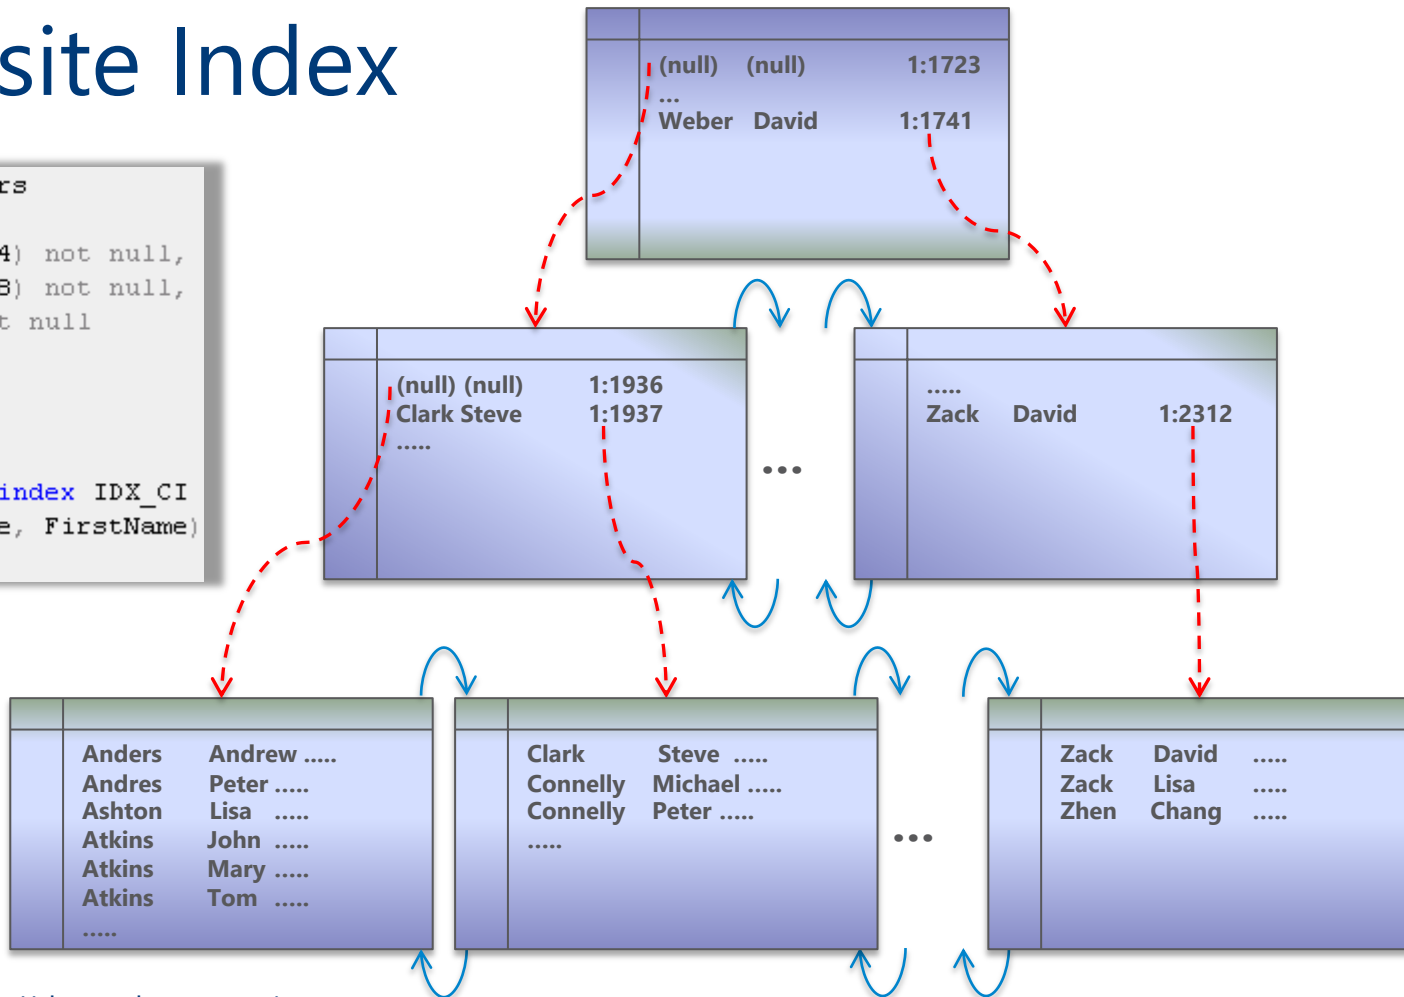
- Unicode parameters for varchar columns (some collations)

### Foreign keys w/o indexes on *referencing* "detail" tables

# Composite Index

```
create table dbo.Customers
(
    FirstName nvarchar(64) not null,
    LastName nvarchar(128) not null,
    Phone varchar(32) not null
    -- ...
)
go

create unique clustered index IDX_CI
on dbo.Customers (LastName, FirstName)
go
```



# Composite Indexes

```
create table dbo.Customers
(
    FirstName nvarchar(64) not null,
    LastName nvarchar(128) not null,
    Phone varchar(32) not null
    -- ...
)
go

create unique clustered index IDX_CI
on dbo.Customers (LastName, FirstName)
go
```

- SARGable when query has SARGable predicates on leftmost column(s) of the index:

where LastName = N'Doe'  
and FirstName = N'John'

where LastName = N'Doe'

- Non-SARGable – all other cases:

where LastName = N'%oe%'  
and FirstName = N'John'

where FirstName = N'John'

# Redundant Indexes

Look at the indexes with the same left-most column(s)

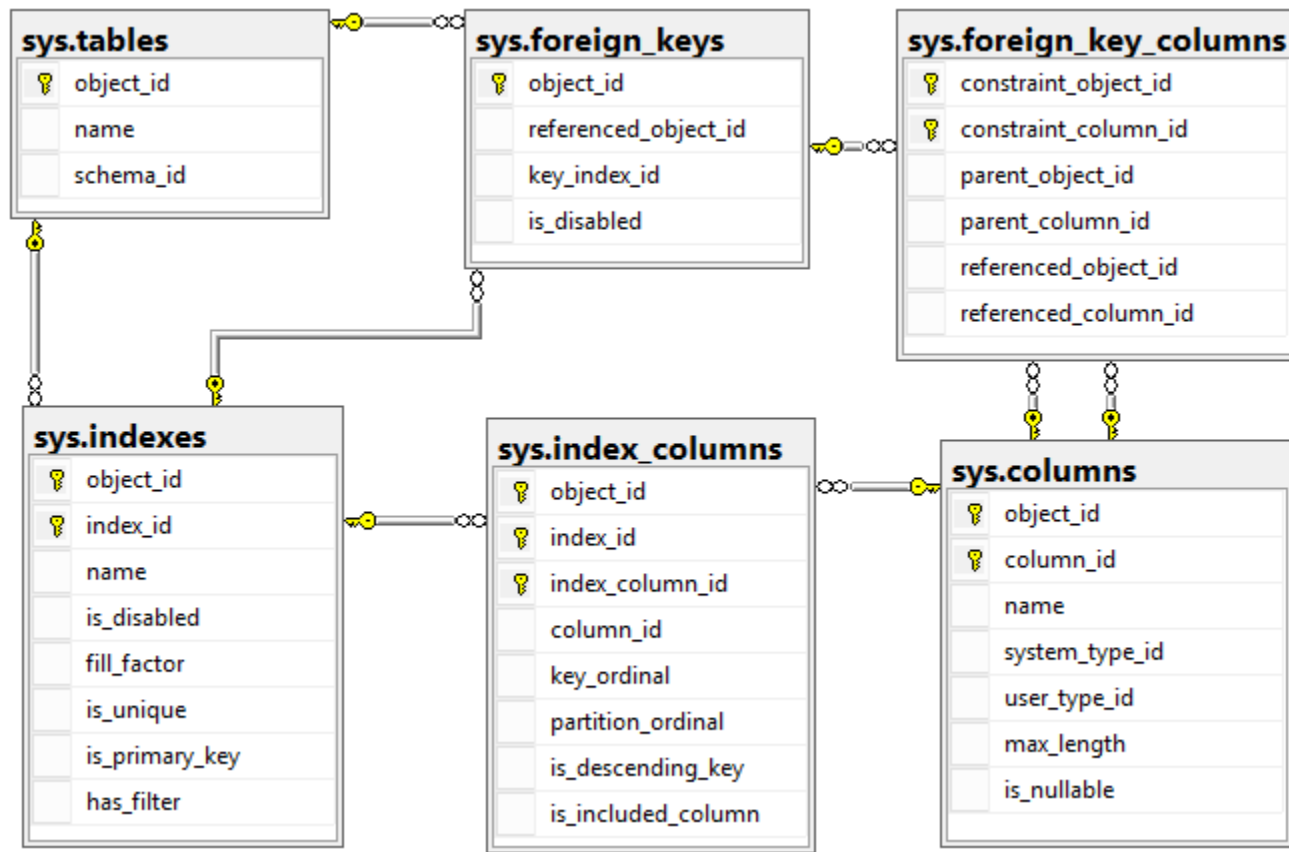
```
create nonclustered index IDX_1  
on dbo.Customers (LastName, FirstName)
```

```
create nonclustered index IDX_2  
on dbo.Customers (LastName)
```

## Multi-Tenant Systems:

- Make *TenantId* as the left-most column
- Factor it to redundant index analysis

# Tales Told By Catalog Views



Heap tables  
Redundant indexes  
Non-indexed FKs  
Wide CI  
Non-unique CI  
Unique identifiers  
Untrusted FKs  
Suboptimal  
FILLFACTOR

Just be creative 😊

# Catalog Views - Ideas

## Heap tables

- `sys.indexes (index_id = 0)`

## Redundant indexes

- Self-join `sys.index_columns (=object_id, >index_id, =column_id, key_ordinal=1)`
- Multi-tenant systems – `key_ordinal > 1`

## Non-indexed foreign keys

- `sys.foreign_key_columns (parent_object_id, parent_column_id, constraint_column_id) IN sys.index_columns(object_id, index_id, key_ordinal)`

## Unique identifiers

- `sys.index_columns (key_ordinal = 1) -> sys.columns.system_type_id = 36`

## Wide clustered indexes

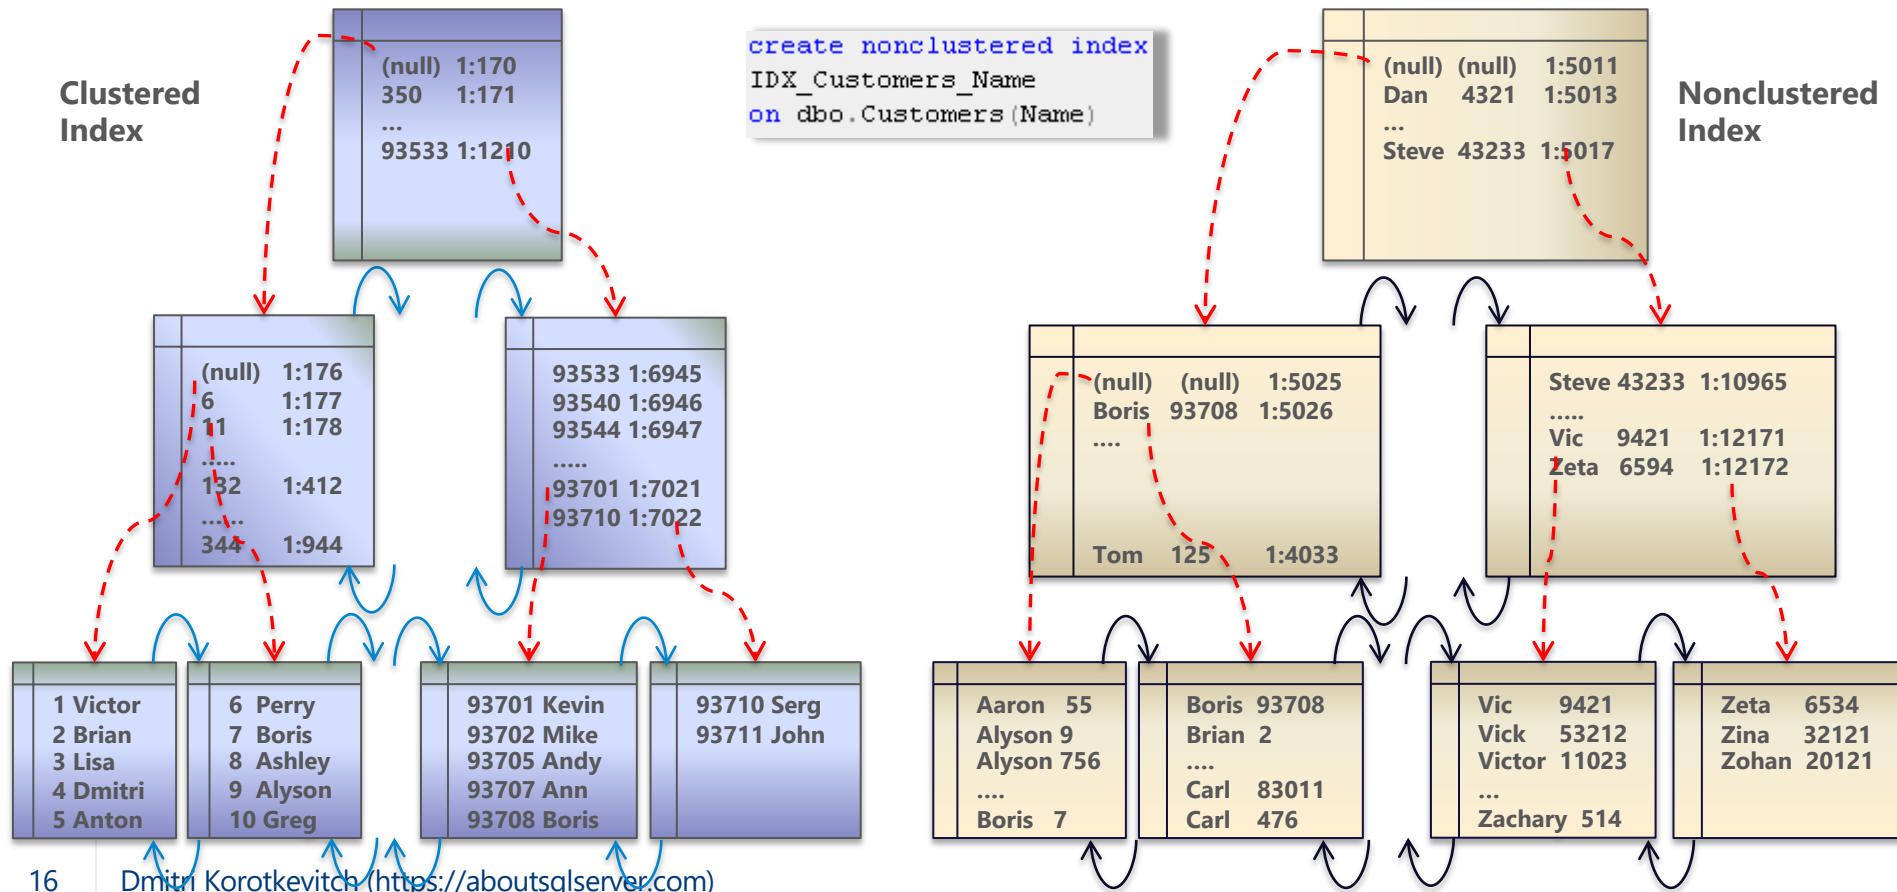
- `sys.indexes (index_id = 1) -> SUM(sys.index_columns -> sys.columns.max_length)`
- `max_length` for (MAX) columns = -1

# Nonclustered Index

Clustered Index

```
create nonclustered index
IDX_Customers_Name
on dbo.Customers (Name)
```

Nonclustered Index



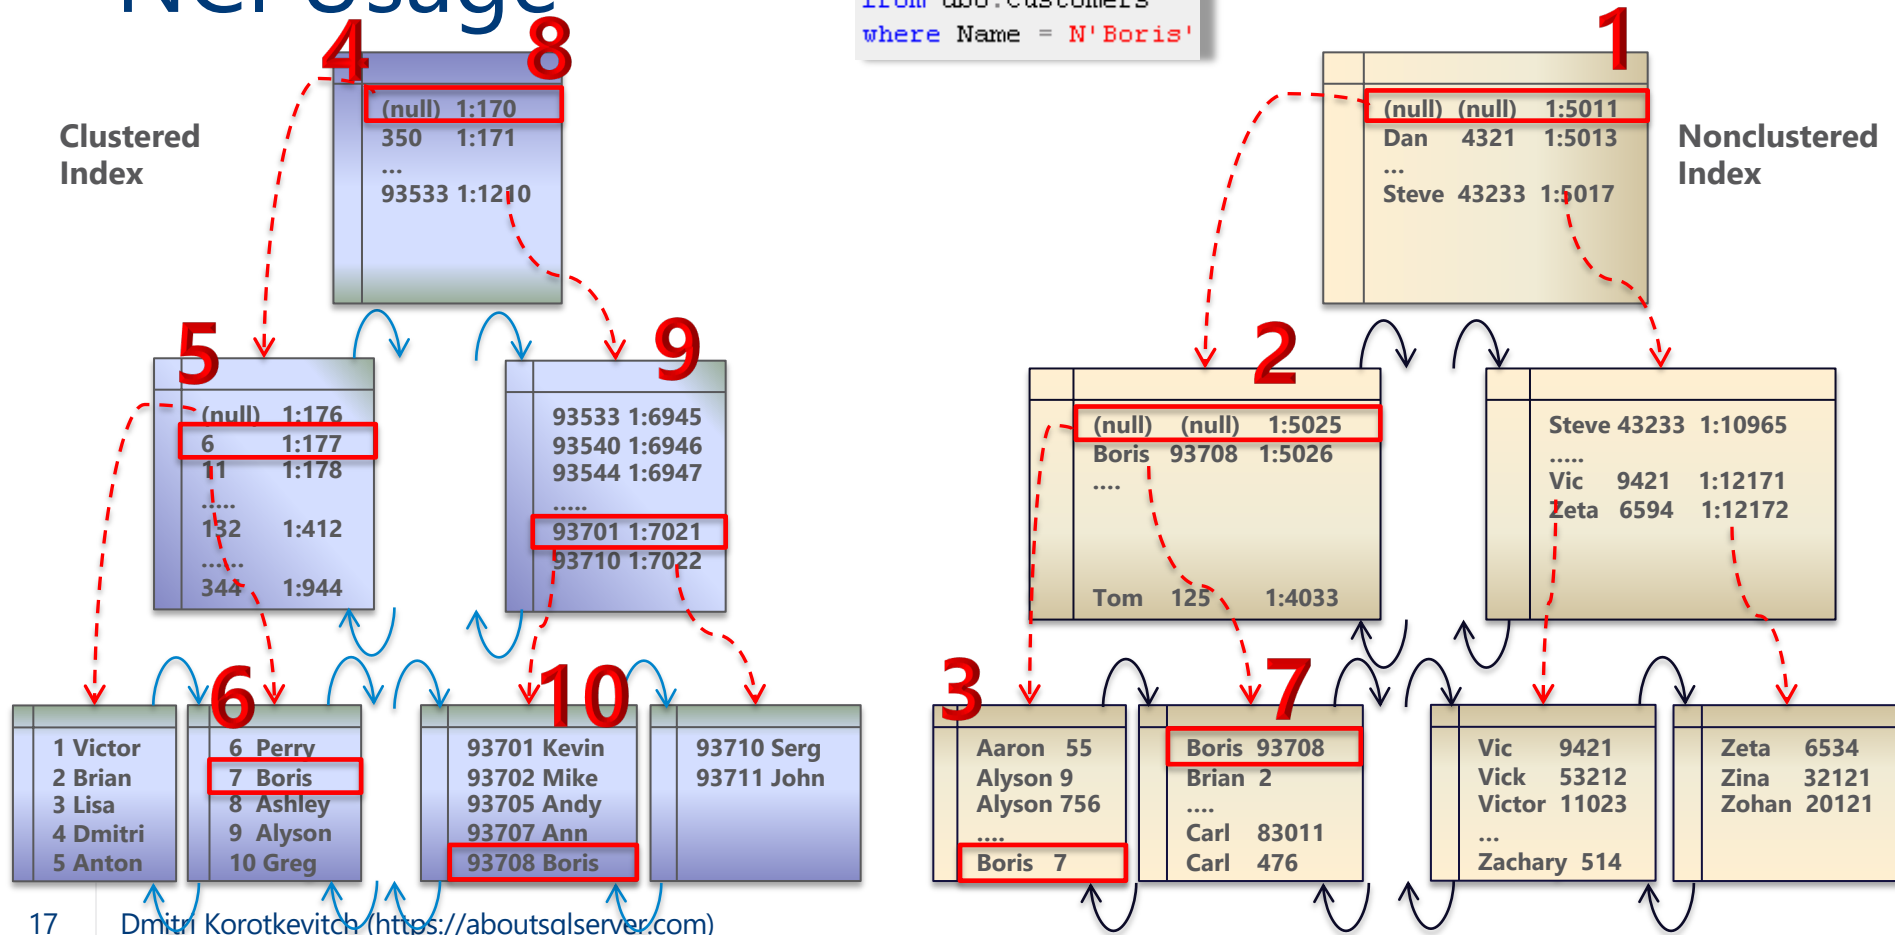


# NCI Usage

```
select *
from dbo.Customers
where Name = N'Boris'
```

Clustered Index

Nonclustered Index





# NCI Usage

## Demo

# sys.dm\_db\_index\_usage\_stats

How often index appears in execution plans

- Seeks: Index Seek operation
- Scans: Index Scan operation
- Lookup: Key Lookup or RID Lookup operations
- Reads: Seeks + Scans + Lookups
- Updates: INSERT + UPDATE + DELETE

	Table	Index	Seeks	Scans	Lookups	Reads	Updates	Last Seek	Last Scan
1	sys.dm_db_index_usage_stats	PK	74567916	0	755160	75323076	1510669	2016-11-18 07:32:15.437	NULL
2	sys.dm_db_index_usage_stats	IX	57940278	0	0	57940278	1510669	2016-11-18 07:32:15.537	NULL
3	sys.dm_db_index_usage_stats	IX	16724	0	0	16724	1510669	2016-11-18 07:00:20.827	NULL
4	sys.dm_db_index_usage_stats	IX	1003195	1090508	0	2093703	1136428	2016-11-18 07:32:08.700	2016-11-18 07:32:08.700
5	sys.dm_db_index_usage_stats	IX	544501	0	0	544501	1510669	2016-11-18 07:32:10.540	NULL
6	sys.dm_db_index_usage_stats	IX	1	0	0	1	680021	2016-11-02 14:17:28.843	NULL
7	sys.dm_db_index_usage_stats	IX	5210970	0	0	5210970	1510669	2016-11-18 07:32:14.423	NULL
8	sys.dm_db_index_usage_stats	IX	0	0	0	0	1510669	NULL	NULL
9	sys.dm_db_index_usage_stats	IX	7049584	0	0	7049584	2403273	2016-11-18 07:32:14.423	NULL
10	sys.dm_db_index_usage_stats	IX	377572	0	0	377572	2403273	2016-11-18 07:32:08.750	NULL

# sys.dm\_db\_index\_operational\_stats

Access methods, I/O, locking, latching activity

- Insert, Update, Delete counts (# of rows)
- singleton\_lookup\_count: Single-row Index Seek operations
- range\_scan: Index Seeks on the range of rows + Index Scans
- LOB and ROW\_OVERFLOW statistics (potential SELECT \* patterns)
- Lock counts and waits on row- and page-levels (good for lock escalation analysis)
- **Page latch count and waits**
- Page IO latch count and waits
- And more..

	index_id	Table	Index	range_scan_count	singleton_lookup_count	row_lock_wait_in_ms	page_latch_wait_in_ms	page_io_latch_wait_in_ms
1	1			1411162	3638897399	0	71286	13634302
2	2			774	0	0	338	283589
3	3			760095	0	0	121	329284
4	4			32726	0	0	828	7878183
5	5			358	0	0	66	3138358
6	6			21	0	0	124	138602
7	8			0	0	0	519	68238
8	33			1	0	0	60	574234
9	34			11012	0	0	490	626016
10	35			0	0	0	126	22540



# Usage and Operational Stats

Demo

# Pattern: Inefficient Reads

High number of Scans in sys.dm\_db\_index\_usage\_stats

	IndexId	Table	Index	Seeks	Scans	Lookups	Updates
1	1			148266434	1500	251782	78441173
2	4			1207312	0	0	48080643
3	5			0	1354122	0	35993284
4	6			73969314	3	0	57187072
5	8			33931425	0	0	35589892
6	3			23019	2134296	0	70957328

Analyze index usefulness (# of seeks, operational overhead)  
Find and optimize queries that use the index



# Analyzing Plan Cache

## Demo

# Pattern: High Maintenance Cost

High Update vs. Read ratio

	IndexId	Table	Index	Seeks	Scans	Lookups	Reads	Updates
1	1			1	9	1121468	1121478	5243411
2	16			8	1	0	9	4940468
3	17			40023	0	0	40023	4940468
4	19			0	0	0	0	5243411
5	23			1089497	0	0	1089497	4940468

Drop unused indexes

Analyze index usefulness (# of seeks, access time, operational overhead)

Find and optimize queries that use the index



# Pattern: Inefficient Clustered Indexes

High Lookup vs. Seek ratio

- Typical pattern: Clustered PRIMARY KEY constraint on identity column

	IndexId	Table	Index	Seeks	Scans	Lookups	Reads	Updates
1	1			1	9	1121468	1121478	5243411
2	16			8	1	0	9	4940468
3	17			40023	0	0	40023	4940468
4	19			0	0	0	0	5243411
5	23			1089497	0	0	1089497	4940468

Change clustered index in the table

# Pattern: Noncovering NCIs

## High Lookup vs. Seek ratio

- Typical pattern: Frequently executed queries use noncovered nonclustered index

	IndexId	Table	Index	Seeks	Scans	Lookups	Reads	Updates
1	1			148266434	1500	251782608	400050542	78441173
2	4			1207312	0	0	1207312	48080643
3	5			988005	1354122	0	988005	35993284
6	35			26734187	0	0	26734187	31211137
7	39			121404098	2	0	121404100	44104235
8	54			9075825	0	0	9075825	31210092
9	57			14180	0	0	14180	7537756
11	59			21019386	5319	0	21024705	48080643
13	61			39117309	6	0	39117315	43700843
14	3			23019	2134296	0	23019039	70957328
15	43			3094705	339090	0	3433795	66895706

Analyze queries with Key Lookups making NCIs covering when possible

# Deeper Dive

## Index size on disk and in buffer pool

- `sys.allocation_units + sys.dm_os_buffer_descriptors` (script is included)

## Lock and latch contention in the system

- `sys.dm_db_index_operational_stats`

## Disk (PAGEIOLATCH) contention

- `sys.dm_db_index_operational_stats + sys.dm_db_index_usage_stats`
  - `page_io_latch_*`
  - Seeks, Scans, singleton\_lookup\_count, range\_scan\_count



# sp\_Index\_Analysis

## Demo

# Caveats

Statistics clears at:

- SQL Server restart
- ALTER INDEX REBUILD in:
  - SQL Server 2012 RTM - SP3 CU2 (Fixed in SP3 CU3)
  - SQL Server 2014 RTM – SP1 (Fixed in CU2)

Remember about  
Readable  
Secondaries!

	Server	IndexId	Table	Index	Seeks	Scans	Lookups	Reads	Updates
1	PRIMARY	1			1	9	1121468	1121478	5243411
2	PRIMARY	16			8	1	0	9	4940468
3	PRIMARY	17			40023	0	0	40023	4940468
4	PRIMARY	19			0	0	0	0	5243411

	Server	IndexId	Table	Index	Seeks	Scans	Lookups	Reads	Updates
1	SECONDARY	1			473	18939	257935	277347	0
2	SECONDARY	16			190645	1337	0	191982	0
3	SECONDARY	17			2860	0	0	2860	0
4	SECONDARY	19			173683	0	0	173683	0

# sys.dm\_db\_index\_physical\_stats

	index_id	partition_number	alloc_unit_type_desc	index_level	page_count	record_count
1	1	1	IN_ROW_DATA	0	651333	12581042
2	1	1	IN_ROW_DATA	1	18746	651333
3	1	1	IN_ROW_DATA	2	202	18746
4	1	1	IN_ROW_DATA	3	1	202
		avg_fragmentation_in_percent	avg_page_space_used_in_percent	version_ghost_record_count	forwarded_record_count	
		5.41	84.98	0	NULL	
		92.78	14.28	0	NULL	
		99.50	42.20	0	NULL	
		0.00	95.71	0	NULL	

avg\_fragmentation\_in\_percent: External fragmentation - # of fragments in DB

avg\_page\_space\_used: Internal fragmentation – page space utilization

Tune FILLFACTOR. Do not use 100 with row versioning (RCSI, Snapshot, etc)

# sys.dm\_db\_index\_physical\_stats

	index_id	partition_number	alloc_unit_type_desc	index_level	page_count	record_count
1	1	1	IN_ROW_DATA	0	651333	12581042
2	1	1	IN_ROW_DATA	1	18746	651333
3	1	1	IN_ROW_DATA	2	202	18746
4	1	1	IN_ROW_DATA	3	1	202
		avg_fragmentation_in_percent	avg_page_space_used_in_percent	version_ghost_record_count	forwarded_record_count	
		5.41	84.98	0	NULL	
		92.78	14.28	0	NULL	
		99.50	42.20	0	NULL	
		0.00	95.71	0	NULL	

Deferred *Ghost* and *Version Store Cleanup* due to row-versioning

Check long running RCSI/Snapshot transaction including readable secondaries

# sys.dm\_db\_index\_physical\_stats

	index_id	partition_number	alloc_unit_type_desc	index_level	page_count	record_count	
1	1	1	IN_ROW_DATA	0	651333	12581042	
2	1	1	IN_ROW_DATA	1	18746	651333	
3	1	1	IN_ROW_DATA	2	202	18746	
4	1	1	IN_ROW_DATA	3	1	202	
		avg_fragmentation_in_percent	avg_page_space_used_in_percent	version_ghost_record_count	forwarded_record_count		
		5.41	84.98	0	NULL		
		92.78	14.28	0	NULL		
		99.50	42.20	0	NULL		
		0.00	95.71	0	NULL		

Inefficient Heap tables (high rate of updates)

Rebuild and consider to remove heap tables



# Key Points



Remember about patterns that prevent *Index Seek*

Use catalog views to identify schema issues

Use usage and operational statistics to identify inefficient indexes

Correlate data from multiple sources during analysis

# Q&A



Blog: <https://aboutsqlserver.com>

E-mail: [dk@aboutsqlserver.com](mailto:dk@aboutsqlserver.com)

Slides / Code: <https://github.com/aboutsqlserver/code>



Thank You